Released_10_07_20

Question 1: Skipped

A mule application is deployed to a cluster of two (2) customer-hosted mule runtimes. Currently, the node named Alice is the primary node and the node named Bob is the secondary node. The Mule application has a flow that polls a directory on a file system for new files.

The primary node Alice fails for an hour and then is restarted.

After the Alice node completely restarts, from what node are the files polled, and what node is now the primary node for the cluster?

○ Files are polled from the Bob node
Alice is now the primary node

○ Files are polled from the Alice node
Alice is now the primary node

○ Files are polled from the Bob node   **(Correct)**
Bob is now the primary node

○ Files are polled from the Alice node
Bob is now the primary node

Question 2: Skipped

In a Mule application, a flow contains two(2) JMS Consume operations that are used to connect to a JMS broker and consume message from two (2) JMS destinations. The Mule application then joins the two consumed JMS message together.

The JMS broker does NOT implement high availability (HA) and periodically experiences scheduled outages of up to 10 minutes for routine maintenance.

What is the most idiomatic (used for its intended purpose) way to build the Mule flow so it can best recover from the expected outages.

○ Configure a transaction for the JMS connector   **(Correct)**

○ Configure a reconnection strategy for the JMS connector

○ Enclose the two (2) JMS operations in a Try scope with an On Error Continue error handler

○ Enclose the two (2) JMS operations in an Until Successful scope

Question 3: Skipped

A mule application is being designed to perform product orchestration. The Mule application needs to join together the responses from an Inventory API and a Product Sales History API with the least latency.

To minimize the overall latency. what is the most idiomatic (used for its intended purpose) design to call each API request in the Mule application?

⭘ Call each API request in a separate Async scope

⭘ Call each API request in a separate lookup call from a DataWeave reduce operator

⭘ Call each API request in a separate route of a Scatter-Gather   **(Correct)**

⭘ Call each API request in a separate route of a Parallel For Each scope

Question 4: Skipped

An automation engineer needs to write script to automate the steps of the API lifecycle, including steps to create, publish, deploy and manage APIs and their implementations in Anypoint Platform.

What Anypoint Platform feature can be used to automate the execution of all these actions in scripts in the easiest way without needing to directly invoke the Anypoint Platform REST APIs?

⭘ The Mule Maven Plugin

⭘ Automated Policies in API Manager

⭘ Runtime Manager agent

⭘ Anypoint CLI   **(Correct)**

Question 5: Skipped

A set of integration mule applications, some of which expose APIs, are being created to enable a new business process. Various stakeholders may be impacted by this. These stakeholders are a combination of semi-technical users (who understand basic integration terminology and concepts such as JSON and XML) and technically skilled potential consumers of the Mule applications and APIs.

What is an effective way for the project team responsible for the mule applications and APIs being built to communicate with these stakeholders using Anypoint platform and its supplied toolset?

○ Use Anypoint design center to implement the mule applications and APIs and give the various stakeholders access to these design center projects, so they can collaborate and provide feedback.

○ Capture documentation about the mule applications and APIs inline within the mule integration flows and use Anypoint Studios. Export documentation feature to provide an HTML version of this documentation to the stakeholders.

○ Use Anypoint exchange to register the various mule applications and APIs and share the RAML definitions with the stakeholders, so they can be discovered.   **(Correct)**

○ Create Anypoint exchange entries with pages elaborating the integration design, including API notebooks (where applicable) to help the stakeholders understand and interact the mule applications and APIs at various levels of technical depth.

Question 6: Skipped

Refer to the exhibits The Mule application is deployed to multiple CloudHub workers with persistent queues enabled.

The retrieveFile flow's event source needs a CSV file from a remote SFTP server and then publishes each record in the CSV file to a VM queue.

The processCustomerRecords flow's VM Listener receives messages from the same VM queue and then processes each message separately.

How are messages routed to the CloudHub workers as messages are received by the VM Listener?
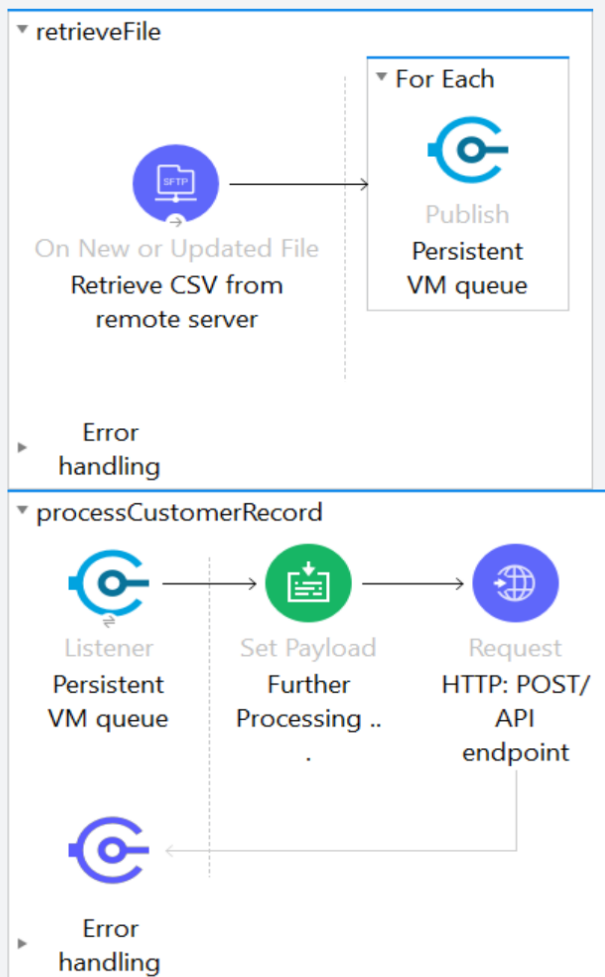
- Each messages is routed to the SAME CloudHub worker that retrieved the file, thereby BINDING ALL messages to ONLY that ONE CloudHub worker **(Correct)**

- Each messages is routed to ONE of the available CloudHub workers in a NON-DETERMINISTIC non-round-robin fashion, thereby APPROXIMATELY BALANCING messages among the CloudHub workers

- Each messages is duplicated to ALL of the CloudHub workers, thereby SHARING EACH message with All the CloudHub workers

- Each message is routed to ONE of the CloudHub workers in a DETERMINISTIC round-robin fashion, thereby EXACTLY BALANCING messages among the CloudHub workers

Question 7: Skipped

An organization has decided on a cloud migration strategy that aims to minimize the organization's own IT resources. Currently, the organization has all of its Mule applications running on its own premises and uses an on-premises load balancer that exposes all APIs under the base URL https://api.acme.com

As part of the migration strategy, the organization plans to migrate all of its Mule applications and load balancer to CloudHub.

What is the most straight-forward and cost-effective approach to Mule application deployment and load balancing that preserves the public URLs?

○ Deploy the Mule applications to CloudHub
Create a CNAME record for api.acme.com in the CloudHub Shared Load Balancer (SLB) pointing to the A record of the on-premises load balancer
Apply mapping rules in the SLB to map URLs to their corresponding Mule applications

○ For each migrated Mule application, deploy an API Proxy Mule application to CloudHub, with all applications under the control of a CloudHub dedicated load balancer (DLB)
Update the CNAME record for api.acme.com in the organization's DNS server pointing to the A record of a CloudHub dedicated load balancer (DLB)
Apply mapping rules in the DLB to map each API Proxy application to its corresponding Mule application

○ Deploy the Mule applications to CloudHub
Update the CNAME record for api.acme.com in the organization's DNS server pointing to the A record of the CloudHub Shared Load Balancer (SLB)
Apply mapping rules in the SLB to map URLs to their corresponding Mule applications

○ Deploy the Mule applications to CloudHub
Update the CNAME record for api.acme.com in the organization's DNS server pointing to the A record of the CloudHub dedicated Load Balancer (DLB)
Apply mapping rules in the DLB to map URLs to their corresponding Mule applications   **(Correct)**

Question 8: Skipped

In Anypoint Platform, a company wants to configure multiple identity providers (IdPs) for various lines of business (LOBs) Multiple business groups and environments have been defined for these LOBs.

What Anypoint Platform feature can use multiple Idps across the company's business groups and environments?

- ○ User management

- ○ Client management   **(Correct)**

- ○ Roles and permissions

- ○ Dedicated load balancers

Question 9: Skipped

A Mule application uses an HTTP Request operation to invoke an external API.

The external API follows the HTTP specification for proper status code usages.

What is a possible cause when a 3xx status code is returned to the HTTP Request operation from the external API?

- ○ The request was NOT RECEIVED by the external API

- ○ The request was REDIRECTED to a different URL by the external API   **(Correct)**

- ○ The request was ACCEPTED by the external API

- ○ The request was NOT ACCEPTED by the external API

Question 10: Skipped

An organization will deploy Mule applications to CloudHub. Business requirements mandate that all application logs be stored ONLY in a external Splunk consolidated logging service and NOT in CloudHub.

In order to most easily store Mule application logs ONLY in Splunk, how must Mule application logging be configured in Runtime Manager, and where should the log4j2 Splunk appender be defined?

○ Disable CloudHub logging in Runtime Manager
Define the Splunk appender in EACH Mule application's log4j2.xml file **(Correct)**

○ Keep the default logging configuration in Runtime Manager
Define the Splunk appender in ONE global log4j2.xml file that is uploaded once to Runtime Manager to support all Mule application deployments

○ Disable CloudHub logging in Runtime Manager
Define the Splunk appender in ONE global log4j2.xml file that is uploaded once to Runtime Manager to support all Mule application deployments

○ Keep the default logging configuration in Runtime Manager
Define the Splunk appender in EACH Mule application's log4j2.xml file

Question 11: Skipped

An organization is migrating all its Mule applications to Runtime Fabric (RTF) None of the Mule application .. projects

Currently, all the Mule applications have been manually deployed to server group among several customer-hosted Mule runtimes Port conflicts between these Mule application deployments are currently managed by the DevOps team who carefully manage Mule application properties files.

When the Mule applications are migrated from the current customer-hosted server group to Runtime Fabric (RTF), do the Mule applications need to be rewritten, and what DevOps port configuration responsibilities change or stay the same?

No, the Mule application do NOT need to be rewritten
DevOps NO LONGER needs to manage port conflicts between the **(Correct)**
Mule applications

YES, the Mule application MUST be rewritten
DevOps NO LONGER needs to manage port conflicts between the Mule
applications

No, the Mule application do NOT need to be rewritten
DevOps MUST STILL manage port conflicts

YES, the Mule application MUST be rewritten
DevOps MUST STILL manage port conflicts

Question 12: Skipped

A Mule application is being designed to receive nightly a CSV file containing millions of records from an external vendor SFTP The records from the file need to be validated, transformed, and then written to a database. Records can be inserted into the database in any order.

In this use case, what combination of Mule components provides the most effective and performant way to write these records to the database?

○ Use a Parallel For Each scope to insert records one by one into the database **(Correct)**

○ Use a Scatter-Gather to bulk insert records into the database

○ Use a Batch Job scope to bulk insert records into the database

○ Use a DataWeave map operation and an Async scope to insert records one by one into the database

Question 13: Skipped

An organization is designing a Mule application to support an all-or-nothing transaction between several database operations and some other connectors so that they all roll back if there is a problem with any of the connectors.

Besides the Database connector, what other connector can be used in the transaction?

○ Anypoint MQ

○ VM **(Correct)**

○ SFTP

○ Object Store

Question 14: Skipped

A Mule application is synchronizing customer data between two different database systems.

What is the main benefits of using XA transactions over local transactions to synchronize these two database systems?

○ Increases throughout

○ Ensures consistency   **(Correct)**

○ Simplifies communication

○ Reduce latency

Question 15: Skipped

A REST API is being designed to implement a Mule application.

What standard interface definition language can be used to define REST APIs?

○ YAML

○ AsyncAPI Specification

○ Web Services Definition Language(WSDL)

○ OpenAPI Specification(OAS)   **(Correct)**

Question 16: Skipped

Anypoint Exchange is required to maintain the source code of some of the assets committed to it, such as connector, Templates and API specifications.

What is the best way to use an organization's source-code management (SCM) system in this context?

Organizations should continue to use an SCM system of their choice, in addition to keeping source code from these asset type in Anypoint Exchange, thereby enabling parallel development, branching and merging.

(Correct)

Organizations need to use Anypoint Exchange as the main SCM system to centralize versioning and avoid code duplication.

Organizations need to point Anypoint Exchange to their SCM system so Anypoint Exchange can pull source code when requested by developers and provide it to Anypoint Studio.

Organizations can continue to use an SCM system of their choice for branching and merging as long as they follow the branching and merging strategy enforced by Anypoint Exchange.

Question 17: Skipped

An organization has previously provisioned its own AWS VPC hosting various servers. The Organization now needs to use CloudHub to host a Mule application that will implement a REST API. Once deployed to CloudHub, this Mule application must be able to communicate securely with the customer-provisioned AWS VPC resources within the same region without being interceptable on the public internet.

What Anypoint Platform features should be used to meet these network communication requirements between CloudHub and the existing customer-provisioned AWS VPC?

○ Add default API Whitelisting policies to API Manager to automatically whitelist the customer-provisioned AWS VPC IP ranges needed by the Mule application

○ Use VM queues in the Mule application to allow any non-Mule assets within the customer-provisioned AWS VPC to subscribe to and receive messages

○ Add a Mulesoft-hosted Anypoint VPC confiured with VPC Peering to the AWS VPC **(Correct)**

○ Configure an external identity provider(Idp) in Anypoint Platform with certificates from the customer-provisioned AWS VPC

Question 18: Skipped

A Mule application named Inventory uses persistent object store. The Inventory Mule application is deployed to CloudHub and is configured to use Object Store v2.

Another Mule application named Roomba is being developed to delete values from the Inventory Mule application's persistent object store. The Rooomba Mule application will also be deployed to CloudHub.

What is the most direct way for the Roomba Mule application to delete values from the Inventory Mule application's persistent object store with the least latency?

○ Use an Anypoint MQ connector configured to directly access the Inventory Mule application's persistent object store

○ Use a VM connector configured to directly access the pesistent queue of the Inventory Mule application's persistent object store

○ Use an Object Store connector configured to access the Inventory Mule application's persistent object store **(Correct)**

○ Use the Object Store v2 REST API configured to access the Inventory Mule applications's persistent object store

Question 19: Skipped

An organization is sizing an Anypoint VPC to extend their internal network to CloudHub.

For this sizing calculation, the organization assumes 150 Mule applications will be deployed among three (3) production environments and will use CloudHub's default zero-downtime feature. Each Mule application is expected to be configured with two (2) CloudHub workers. This is expected to result in several Mule application deployments per hour.

What is the minimum number of IP addresses that should be configured for this VPC that results in the smallest usable private IP address range to support the deployment and zero downtime of these 150 Mule applications (not accounting for any future Mule applications)?

- ○ 10.0.0/21 (2048 IPs)   **(Correct)**

- ○ 10.0.0/22 (1024 IPs)

- ○ 10.0.0/23 (512 IPs)

- ○ 10.0.0/24 (256 IPs)

Question 20: Skipped

An organization uses a four(4) node customer-hosted Mule runtime cluster to host one(1) stateless API implementation. The API is accessed over HTTPS through a load balancer that uses round-robin for load distribution. Each node in the cluster has been sized to able to accept four(4) times the current number of requests.

Two(2) nodes in the cluster experience a power outage and are no longer available. The load balancer detects the outage and blocks the two unavailable nodes from receiving further HTTP requests.

What performance-related consequence is guaranteed to happen, on average, assuming the remaining cluster nodes are fully operational?

○ 50% reduction in the throughout of the API

○ 100% increase in the number of the requests received by each remaining node

○ 100% increase in the average response time of the API

○ 50% increase in the JVM heap memory consumed by each remaining node   **(Correct)**

Question 21: Skipped

A Mule application is deployed to a single CloudHub worker and the public URL appears in Runtime Manager as the App url.

Requests are sent by external web clients over the public internet to the Mule application's App url. Each of these requests is routed to the HTTPS Listener event source of the running Mule application.

Later, the DevOps team edits some properties of this running Mule application in Runtime Manager.

Immediately after the new property values are applied in Runtime Manager, how is the current Mule application deployment affected, and how will future web client requests to the Mule application be handled?

CloudHub will redeploy the Mule application to a NEW CloudHub worker
New web client requests are ROUTED to the OLD CloudHub worker until the NEW CloudHub worker is available **(Correct)**

CloudHub will redeploy the Mule application to a NEW CloudHub worker
New web client requests will RETURN AN ERROR until the NEW CloudHub worker is available

CloudHub will redeploy the Mule application to the OLD Cloud Hub worker
New web client requests are ROUTED to the OLD CloudHub worker BOTH before and after the Mule application is redeployed

CloudHub will redeploy the Mule application to the OLD CloudHub worker
New web client requests will RETURN AN ERROR until the Mule application is redeployed to the OLD CLoudHub worker

Question 22: Skipped

An organization is evaluating using the CloudHub Shared Load Balancer (SLB) vs. creating a CloudHub dedicated load balancer (DLB). They are evaluating how this choice affects the various types of certificates used by CloudHub-deployed Mule applications, including MuleSoft-provided, customer-provided, or Mule applicaton-provided certificates.

What type of restrictions exist on the types of certificates that can be exposed by the CloudHub Shared Load Balancer (SLB) to external web clients over the public internet?

○ Only customer-provided self-signed certificates are exposed

○ Only underlying Mule application certificates are exposed(pass-through)

○ Only MuleSoft-provided certificates are exposed   **(Correct)**

○ Only customer-provided wildcard certificates are exposed